

November 15, 2023 Class 13: Transcriptomics and the analysis of RNA-Seq data

Savannah Bogus A69027475

```
counts<-read.csv("airway_scaledcounts.csv",row.names=1)
metadata<-read.csv("airway_metadata.csv")
```

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG000000000419	467	523	616	371	582
ENSG000000000457	347	258	364	237	318
ENSG000000000460	96	81	73	66	118
ENSG000000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG000000000419	781	417	509		
ENSG000000000457	447	330	324		
ENSG000000000460	94	102	74		
ENSG000000000938	0	0	0		

```
head(metadata)
```

	id	dex	celltype	geo_id
1	SRR1039508	control	N61311	GSM1275862
2	SRR1039509	treated	N61311	GSM1275863
3	SRR1039512	control	N052611	GSM1275866
4	SRR1039513	treated	N052611	GSM1275867

```
5 SRR1039516 control N080611 GSM1275870
6 SRR1039517 treated N080611 GSM1275871
```

Q1

How many genes are in the dataset? 38,694 genes.

Q2

How many control cell lines are there?

```
table(metadata$dex)
```

```
control treated
      4       4
```

4 control cell lines.

Toy differential gene expression

I want to compare the treated and control columns. In order to do this, we have to 1. identify/separate out control columns 2. calculate mean value per gene, save as “control.mean” 3. repeat for treated 4. compare

```
control <- metadata[metadata[, "dex"]=="control",]
control.counts <- counts[, control$id]
control.mean <- rowSums( control.counts )/4
head(control.mean)
```

```
ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
      900.75           0.00           520.50           339.75           97.25
ENSG0000000000938
      0.75
```

Or we can use tidyverse

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

```
filter, lag
```

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

```
control <- metadata %>% filter(dex=="control")
control.counts <- counts %>% select(control$id)
control.mean <- rowSums(control.counts)/4
head(control.mean)
```

```
ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460
          900.75           0.00           520.50           339.75           97.25
ENSG000000000938
          0.75
```

I personally think I like tidyverse better because I really like that %>% filter function.

Q3

How could you add a function that would help? The code above only works if you know how many control values you have because it is hard coded in.

Q4

```
treated <- metadata %>% filter(dex=="treated")
treated.counts <- counts %>% select(treated$id)
treated.mean <- rowMeans(treated.counts)
head(treated.mean)
```

ENSG000000000003	ENSG000000000005	ENSG000000000419	ENSG000000000457	ENSG000000000460
658.00	0.00	546.00	316.50	78.75
ENSG000000000938				
0.00				

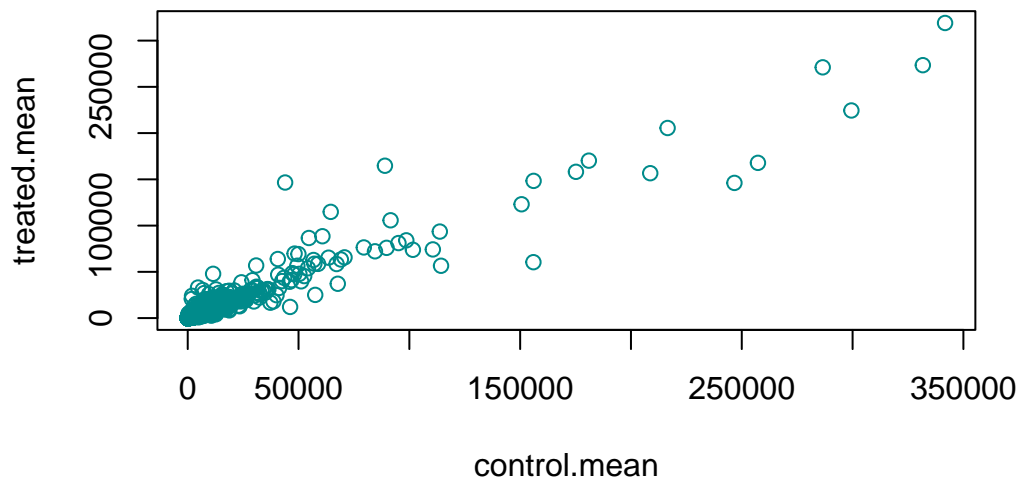
```
mean.counts<-data.frame(control.mean, treated.mean)
colSums(mean.counts)
```

```
control.mean treated.mean
      23005324      22196524
```

Q5a


plot of treated samples vs control samples

```
plot(control.mean,treated.mean, col=c("darkcyan"))
```



Now I'm gonna use ggplot :)

```
library(ggplot2)
ggplot(mean.counts)+aes(treated.mean,control.mean)+
  geom_point(color="darkcyan")
```




There's like 60k data points, but I cannot hardly see them, so I'll use a log scale to try to see them.

```
plot(treated.mean,control.mean,log="xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 y values <= 0 omitted from logarithmic plot



Now, we're going to add log2 info because it tends to have better mathematical properties. And then we'll add it to the mean.counts data.frame.

```
mean.counts$log2fc<-log2(mean.counts$treated.mean/  
                        mean.counts$control.mean)  
head(mean.counts)
```

	control.mean	treated.mean	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000938	0.75	0.00	-Inf

There's several examples with no expression (ie. the NaN, which result from trying to divide by 0, and the -Inf, which results from trying to take the log of a 0.). We should remove those.

```
zero.values<- which(mean.counts[,1:2]==0, arr.ind=TRUE)  
  
to.rm<-unique(zero.values[,1])
```

```
mycounts<-mean.counts[-to.rm,]
```

#Q7

The purpose of the `arr.ind` I didn't understand initially, so I had to ask Claude. It is used to extract elements from arrays, and without that portion of the code, `zero.values` is listed as values, and with that bit of code, it is listed as data and is something I can actually click on and look at. And `unique` is going to keep us from double counting rows that have 0 in multiple values.

Q9 and Q10

Next, we're going to do up and down regulation.

```
nrow(mycounts)
```

```
[1] 21817
```

```
upreg<-mycounts$log2fc>2  
downreg<-mycounts$log2fc< (-2)
```

```
sum(upreg)
```

```
[1] 250
```

```
sum(downreg)
```

```
[1] 367
```

There are 21817 genes left that did not have a 0, and there are 250 upregulated and 367 downregulated genes.

Q11

Do I trust these results? We haven't look at any p-values yet, which are usually a huge portion of volcano plots.

Setting up for DESeq

```
library(DESeq2)
```

Loading required package: S4Vectors

Loading required package: stats4

Loading required package: BiocGenerics

Attaching package: 'BiocGenerics'

The following objects are masked from 'package:dplyr':

combine, intersect, setdiff, union

The following objects are masked from 'package:stats':

IQR, mad, sd, var, xtabs

The following objects are masked from 'package:base':

anyDuplicated, aperm, append, as.data.frame, basename, cbind,
colnames, dirname, do.call, duplicated, eval, evalq, Filter, Find,
get, grep, grepl, intersect, is.unsorted, lapply, Map, mapply,
match, mget, order, paste, pmax, pmax.int, pmin, pmin.int,
Position, rank, rbind, Reduce, rownames, sapply, setdiff, sort,
table, tapply, union, unique, unsplit, which.max, which.min

Attaching package: 'S4Vectors'

The following objects are masked from 'package:dplyr':

first, rename

The following object is masked from 'package:utils':

findMatches

The following objects are masked from 'package:base':

expand.grid, I, unname

Loading required package: IRanges

Attaching package: 'IRanges'

The following objects are masked from 'package:dplyr':

collapse, desc, slice

Loading required package: GenomicRanges

Loading required package: GenomeInfoDb

Warning: package 'GenomeInfoDb' was built under R version 4.3.2

Loading required package: SummarizedExperiment

Warning: package 'SummarizedExperiment' was built under R version 4.3.2

Loading required package: MatrixGenerics

Loading required package: matrixStats

Attaching package: 'matrixStats'

The following object is masked from 'package:dplyr':

count

Attaching package: 'MatrixGenerics'

The following objects are masked from 'package:matrixStats':

colAlls, colAnyNAs, colAnys, colAvgPerRowSet, colCollapse,
colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
colWeightedMeans, colWeightedMedians, colWeightedSds,
colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgPerColSet,
rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
rowWeightedSds, rowWeightedVars

Loading required package: Biobase

Welcome to Bioconductor

Vignettes contain introductory material; view with
'browseVignettes()'. To cite Bioconductor, see
'citation("Biobase")', and for packages 'citation("pkgname")'.

Attaching package: 'Biobase'

The following object is masked from 'package:MatrixGenerics':

rowMedians

The following objects are masked from 'package:matrixStats':

anyMissing, rowMedians

```
dds<-DESeqDataSetFromMatrix(countData=counts,
                             colData=metadata,
                             design=~dex)
```

converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in design formula are characters, converting to factors

```
dds<-DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

To get results from this dds thing in a usable way, use the DESeq results() function

```
res<-results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
ENSG000000000005	0.000000	NA	NA	NA	NA
ENSG000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
ENSG000000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106


ENSG00000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
ENSG00000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029
	padj				
	<numeric>				
ENSG00000000003	0.163035				
ENSG00000000005	NA				
ENSG00000000419	0.176032				
ENSG00000000457	0.961694				
ENSG00000000460	0.815849				
ENSG00000000938	NA				

Data Visualization

Do you remember the volcano plot from the paper discussion today? We're gonna make one.

```
mycolors<- rep("cyan",nrow(res))
mycolors[res$log2FoldChange>2]<-"darkcyan"
mycolors[res$log2FoldChange<(-2)]<-"darkcyan"
mycolors[res$padj > 0.05]<-"darkorange3"

plot(res$log2FoldChange,-log(res$padj),
      xlab="-log P-value",
      ylab="Log2 Fold-Change",
      col=mycolors,
      abline(v=c(2,-2),col="violet",lty=2)
)
abline(h=-log(0.05), col="pink", lty=2)
```



```
write.csv(res,file="myresults.csv")
```

Adding Annotation Data

We need to translate/map our ensemble IDs into gene names or else we don't know what the f is going on.

```
library(AnnotationDbi)
```

Attaching package: 'AnnotationDbi'

The following object is masked from 'package:dplyr':

```
select
```

```
library("org.Hs.eg.db")
```

`org.Hs.eg.db` is in a special format, so we have to use a special function `columns(org.Hs.eg.db)` in order to read it.

```
columns(org.Hs.eg.db)
```

```
[1] "ACCNUM"      "ALIAS"      "ENSEMBL"    "ENSEMBLPROT" "ENSEMBLTRANS"
[6] "ENTREZID"    "ENZYME"     "EVIDENCE"   "EVIDENCEALL" "GENENAME"
[11] "GENETYPE"    "GO"         "GOALL"      "IPI"         "MAP"
[16] "OMIM"        "ONTOLOGY"   "ONTOLOGYALL" "PATH"        "PFAM"
[21] "PMID"        "PROSITE"    "REFSEQ"     "SYMBOL"      "UCSCKG"
[26] "UNIPROT"
```

Our current data uses ENSEMBL IDs, and we're gonna map to SYMBOL, via `mapIds()`

```
res$symbol <- mapIds(org.Hs.eg.db,
                     keys=row.names(res), # Our genenames
                     keytype="ENSEMBL",  # The format of our genenames
                     column="SYMBOL",     # The new format we want to add
                     multiVals="first") #this means what do we do if there's multiple values
```

'select()' returned 1:many mapping between keys and columns

```
head(res$symbol)
```

```
ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
      "TSPAN6"      "TNMD"      "DPM1"      "SCYL3"      "FIRRM"
ENSG0000000000938
      "FGR"
```

We're going to add a few more mappings because we want to look at pathways.

Q11

We're gonna run the `mapIds` 2x more to add Entrez ID and UniProt accession

```
res$entrez<-mapIds(org.Hs.eg.db,
                  keys=row.names(res),
```

```
keytype="ENSEMBL",
column="ENTREZID",
multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
res$uniprot<-mapIds(org.Hs.eg.db,
  keys=row.names(res), # Our genenames
  keytype="ENSEMBL",   # The format of our genenames
  column="UNIPROT",     # The new format we want to add
  multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
res$genename<-mapIds(org.Hs.eg.db,
  keys=row.names(res),
  keytype="ENSEMBL",
  column="GENENAME",
  multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

We can order results by adjusted p value

```
ordered<-order(res$padj)

head(res[ordered,])
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 10 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG00000152583	954.771	4.36836	0.2371268	18.4220	8.74490e-76
ENSG00000179094	743.253	2.86389	0.1755693	16.3120	8.10784e-60
ENSG00000116584	2277.913	-1.03470	0.0650984	-15.8944	6.92855e-57
ENSG00000189221	2383.754	3.34154	0.2124058	15.7319	9.14433e-56
ENSG00000120129	3440.704	2.96521	0.2036951	14.5571	5.26424e-48

	padj	symbol	entrez	uniprot
	<numeric>	<character>	<character>	<character>
ENSG00000148175	13493.920			
		SPARCL1	8404	AOA024RDE1
ENSG00000179094	6.13966e-56	PER1	5187	015534
ENSG00000116584	3.49776e-53	ARHGEF2	9181	Q92974
ENSG00000189221	3.46227e-52	MAOA	4128	P21397
ENSG00000120129	1.59454e-44	DUSP1	1843	B4DU40
ENSG00000148175	1.83034e-42	STOM	2040	F8VSL7
		genename		
		<character>		
ENSG00000152583		SPARC like 1		
ENSG00000179094		period circadian reg..		
ENSG00000116584		Rho/Rac guanine nucl..		
ENSG00000189221		monoamine oxidase A		
ENSG00000120129		dual specificity pho..		
ENSG00000148175		stomatin		

Pathway Analysis

To do this, we're gonna install a few more packages.

```
library(pathview)
```

```
#####
Pathview is an open source software package distributed under GNU General
Public License version 3 (GPLv3). Details of GPLv3 is available at
http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to
formally cite the original Pathview paper (not just mention it) in publications
or products. For details, do citation("pathview") within R.
```

The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG license agreement (details at <http://www.kegg.jp/kegg/legal.html>).

```
#####
```

```
library(gage)
```



```
library(gageData)
```

```
data(kegg.sets.hs)  
head(kegg.sets.hs,2)
```

```
$`hsa00232 Caffeine metabolism`
```

```
[1] "10" "1544" "1548" "1549" "1553" "7498" "9"
```

```
$`hsa00983 Drug metabolism - other enzymes`
```

```
[1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"  
[9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"  
[17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"  
[25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"  
[33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"  
[41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"  
[49] "8824" "8833" "9" "978"
```

gage needs a vector it won't know what to do with DESeq stuff, and gage speaks ENTREZ, so names have to be in ENTREZ format

```
foldchanges=res$log2FoldChange  
names(foldchanges)=res$entrez  
head(foldchanges)
```

```
          7105          64102          8813          57147          55732          2268  
-0.35070302          NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

```
keggres=gage(foldchanges,gsets=kegg.sets.hs)  
attributes(keggres)
```

```
$names
```

```
[1] "greater" "less" "stats"
```

Next, look at the first 3 less than genes.

```
head(keggres$less,3)
```

hsa05332 is the kegg identifier. We're gonna look at asthma because that's what he used to work on. So now we'll make a pathway viewer.

'select()' returned 1:1 mapping between keys and columns

Info: Writing image file hsa05310.pathview.png

